

```

# This Python 3 environment comes with many helpful analytics
libraries installed
# It is defined by the kaggle/python Docker image:
https://github.com/kaggle/docker-python
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/"
directory
# For example, running this (by clicking run or pressing Shift+Enter)
will list all files under the input directory

import os
for dirname, _, filenames in os.walk('/kaggle/input'):
    for filename in filenames:
        print(os.path.join(dirname, filename))

# You can write up to 20GB to the current directory (/kaggle/working/)
that gets preserved as output when you create a version using "Save &
Run All"
# You can also write temporary files to /kaggle/temp/, but they won't
be saved outside of the current session

import io
import pickle
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

tdf =
pickle.load(open('../input/testandtrain/test_with_lemm_2.pkl','rb'))
main_df =
pickle.load(open('../input/testandtrain/train_with_lemm_2.pkl','rb'))

tdf['text']

0          interview execution
1      know post truthful admins hate wont anything e...
2          z youre real bore go bore someone else twt
3      yet remain confusion mention sac withdraw berg...
4          refer lose interest vandalise talk page dark
...
38292      please refrain remove content wikipedia storm ...
38293                                          utc
38294      list girl alphabetical order kep move victoria...
38295      dumb fuck delete angry nintendo nerd page dumb...
38296      fine job would never find log lok carefully ima...
Name: text, Length: 38297, dtype: object

```

main\_df

	id \
0	a8be7c5d4527adbbf15f
1	0b7ca73f388222aad64d
2	db934381501872ba6f38
3	228015c4a87c4b1f09a7
4	b18f26cfa1408b52e949
...	...
89354	748a13233clea91c4584
89355	e49b832cc766ee220113
89356	ff4751b348157ac2b585
89357	94a3215b11a49ba57d61
89358	fa56881999d000a21cf1

	text	harsh \
0	utci interest argue policy resolve ongoing con...	0
1	aded mising parameter templateinfobox organiza...	0
2	sandbox madre sandbox	1
3	god sir sir obviously comprehend importance sc...	1
4	source incase forget someone else want pick gr...	0
...	...	...
89354	becuase critic actualy read boks	0
89355	youre go technical boyd never post goglegroups...	0
89356	join u fb helo pakistani wikipedians asalamoal...	0
89357	lockhart update reference denis lockhart mwala	0
89358	whole section bad format neds writen clear	0

	extremely_harsh	vulgar	threatening	disrespect	targeted_hate
0	0	0	0	0	0
1	0	0	0	0	0
2	0	0	0	0	0
3	0	1	1	1	0
4	0	0	0	0	0
...	...	...	...	...	...
89354	0	0	0	0	0
89355	0	0	0	0	0
89356	0	0	0	0	0
89357	0	0	0	0	0

```
89358          0          0          0          0          0
```

```
[89359 rows x 8 columns]
```

```
tdf.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 38297 entries, 0 to 38296
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype
---  -
 0    id      38297 non-null    object
 1   text    38297 non-null    object
dtypes: object(2)
memory usage: 598.5+ KB
```

## Classifier Chain

```
# using classifier chains
```

```
from skmultilearn.problem_transform import ClassifierChain
from sklearn.linear_model import LogisticRegression
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
x_train = main_df['text'].head(15000)
x_train
```

```
0      utci interest argue policy resolve ongoing con...
1      aded mising parameter templateinfo box organiza...
2                                     sandbox madre sandbox
3      god sir sir obviously comprehend importance sc...
4      source incase forget someone else want pick gr...
```

```
...
14995  thank im expert string instrument help bagpipe...
14996  actually mention source late article open artic...
14997  rofl thats thing hasnt hapened yet along selfb...
14998  age since time doesnt realy pa show homer age ...
14999  like great job thanksi especialy like fact peo...
```

```
Name: text, Length: 15000, dtype: object
```

```
y_train = main_df.drop(labels = ['id', 'text'], axis=1).head(15000)
y_train
```

```
      harsh  extremely_harsh  vulgar  threatening  disrespect
targeted_hate
0          0                0        0            0            0
0
1          0                0        0            0            0
0
2          1                0        0            0            0
0
```

3	1	0	1	1	1
0					
4	0	0	0	0	0
0					
...	...	...	...	...	...
...					
14995	0	0	0	0	0
0					
14996	0	0	0	0	0
0					
14997	0	0	0	0	0
0					
14998	0	0	0	0	0
0					
14999	0	0	0	0	0
0					

[15000 rows x 6 columns]

```
x_test = tdf['text']
type(x_test)
```

pandas.core.series.Series

```
tfidf = TfidfVectorizer()
```

```
x_train_tfidf = tfidf.fit_transform(x_train)
x_test_tfidf = tfidf.transform(x_test)
```

```
# initialize classifier chains multi-label classifier
```

```
classifier = ClassifierChain(LogisticRegression())
```

```
# Training logistic regression model on train data
```

```
classifier.fit(x_train_tfidf, y_train)
```

```
ClassifierChain(classifier=LogisticRegression(), require_dense=[True,
True])
```

```
# predict
```

```
y_pred = classifier.predict_proba(x_test_tfidf)
```

```
y_pred
```

```
<38297x6 sparse matrix of type '<class 'numpy.float64'>'
with 229782 stored elements in Compressed Sparse Column format>
```

```
labels = main_df.iloc[:,2:8]
```

```
labels
```

	harsh	extremely_harsh	vulgar	threatening	disrespect
targeted_hate					
0	0	0	0	0	0
0					

1	0	0	0	0	0
0					
2	1	0	0	0	0
0					
3	1	0	1	1	1
0					
4	0	0	0	0	0
0					
...	...	...	...	...	...
...					
89354	0	0	0	0	0
0					
89355	0	0	0	0	0
0					
89356	0	0	0	0	0
0					
89357	0	0	0	0	0
0					
89358	0	0	0	0	0
0					

[89359 rows x 6 columns]

y\_pred = y\_pred.toarray()

y\_pred\_df = pd.DataFrame(y\_pred, columns = labels.columns)

y\_pred\_df

	harsh	extremely_harsh	vulgar	threatening	disrespect \
0	0.104800	0.000434	0.003618	0.000612	0.003690
1	0.223649	0.000247	0.004230	0.000542	0.003266
2	0.260670	0.000412	0.005378	0.000697	0.004370
3	0.051655	0.000377	0.004041	0.000593	0.002871
4	0.044735	0.000317	0.002791	0.000507	0.002061
...	...	...	...	...	...
38292	0.006338	0.000229	0.002424	0.000462	0.001774
38293	0.018515	0.000413	0.002908	0.000582	0.003344
38294	0.039060	0.000395	0.004204	0.000546	0.003274
38295	0.986294	0.234234	0.931438	0.025603	0.905458
38296	0.020765	0.000277	0.002789	0.000577	0.002523

	targeted_hate
0	0.000846
1	0.000702
2	0.000805
3	0.000711
4	0.000661
...	...
38292	0.000477
38293	0.000778
38294	0.000604

```
38295      0.088567
38296      0.000762
```

```
[38297 rows x 6 columns]
```

```
submission_df = pd.concat([tdf.id,y_pred_df],axis=1)
submission_df.to_csv('submission.csv',index = False)
```

## Label Powerset

```
# using Label Powerset
```

```
from skmultilearn.problem_transform import LabelPowerset
from sklearn.linear_model import LogisticRegression
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
x_train = main_df['text'].head(10000)
x_train
```

```
y_train = main_df.drop(labels = ['id','text'], axis=1).head(10000)
y_train
```

```
x_test = tdf['text']
x_test
```

```
tfidf = TfidfVectorizer()
```

```
x_train_tfidf = tfidf.fit_transform(x_train)
x_test_tfidf = tfidf.transform(x_test)
```

```
classifier = LabelPowerset(LogisticRegression())
# Training logistic regression model on train data
classifier.fit(x_train_tfidf, y_train)
```

```
# predict
```

```
y_pred = classifier.predict_proba(x_test_tfidf)
```

```
y_pred = y_pred.toarray()
```

```
labels = main_df.iloc[:,2:8]
labels
```

```
y_pred.shape
```

```
y_pred_df = pd.DataFrame(y_pred,columns = labels.columns)
y_pred_df
```

```
submission_df = pd.concat([tdf.id,y_pred_df],axis=1)
submission_df.to_csv('submission.csv',index = False)
```

## MLknn

```
#from skmultilearn.adapt import MLkNN
```

```
from scipy.sparse import csr_matrix, lil_matrix
from sklearn.feature_extraction.text import TfidfVectorizer
```

```

#k=6 since there are 6 labels
classifier_new = MLkNN(k=6)

x_train = main_df['text'].head(10000)
x_train

y_train = main_df.drop(labels = ['id', 'text'], axis=1).head(10000)
y_train

x_test = tdf['text']
x_test

tfidf = TfidfVectorizer()

x_train_tfidf = tfidf.fit_transform(x_train)
x_test_tfidf = tfidf.transform(x_test)

x_train_tfidf = lil_matrix(x_train_tfidf).toarray()
y_train = lil_matrix(y_train).toarray()
x_test_tfidf = lil_matrix(x_test_tfidf).toarray()

# train
classifier_new.fit(x_train_tfidf, y_train)

# predict
y_pred = classifier_new.predict(x_test_tfidf)

```

## Linear SVC

```

from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.svm import LinearSVC
from sklearn.multiclass import OneVsRestClassifier
from sklearn.pipeline import Pipeline
from sklearn.calibration import CalibratedClassifierCV
from sklearn.metrics import classification_report

X_train = main_df['text']

y_train = main_df.drop(labels = ['id', 'text'], axis=1)

categories = main_df.iloc[:,2:8].columns

X_test = tdf['text']

SVC_pipeline = Pipeline([
    ('tfidf', TfidfVectorizer()),
    ('clf', OneVsRestClassifier(LinearSVC()))])

for category in categories:
    print('... Processing {}'.format(category))
    # train the model using X_dtm & y
    SVC_pipeline.fit(X_train, y_train[category])
    # compute the testing accuracy
    prediction = SVC_pipeline.predict(X_test)

```

since we have to find probability of labels we cannot directly use onevsrest because of the error : This 'OneVsRestClassifier' has no attribute 'predict\_proba'

```
SVC_pipeline = Pipeline([
    ('tfidf', TfidfVectorizer()),
    ('clf', CalibratedClassifierCV(LinearSVC(C = 10))),
])
cla = OneVsRestClassifier(SVC_pipeline)
cla.fit(X_train, y_train)
```

```
OneVsRestClassifier(estimator=Pipeline(steps=[('tfidf',
                                                TfidfVectorizer()),
                                                ('clf',
```

```
CalibratedClassifierCV(base_estimator=LinearSVC(C=10))]))
```

```
y_pred_svc = cla.predict_proba(X_test)
```

```
y_pred_svc.shape
```

```
(38297, 6)
```

```
y_pred_svc
```

```
array([[1.70391107e-01, 4.44953455e-03, 1.97591234e-03, 6.81200640e-
04,
        1.79114300e-03, 1.96399698e-03],
       [5.60333981e-01, 1.06906238e-02, 2.79940220e-02, 2.02496465e-
03,
        1.06263908e-01, 3.09471682e-03],
       [2.92418769e-01, 3.10111041e-03, 1.02386010e-01, 1.06203947e-
03,
        6.31791555e-02, 7.71424120e-03],
       ...,
       [6.97491675e-03, 5.16868318e-04, 8.57713607e-03, 1.77557770e-
04,
        2.28373323e-02, 1.17974089e-03],
       [9.9998869e-01, 5.78736406e-01, 9.99903357e-01, 8.99214208e-
03,
        9.77605332e-01, 2.53610878e-02],
       [9.98316264e-03, 2.28602266e-03, 1.20480085e-03, 1.46304882e-
04,
        2.38564816e-03, 8.47497354e-04]])
```

```
y_pred_df = pd.DataFrame(y_pred_svc, columns = categories)
y_pred_df
```

	harsh	extremely_harsh	vulgar	threatening	disrespect	\
0	0.170391	0.004450	0.001976	0.000681	0.001791	
1	0.560334	0.010691	0.027994	0.002025	0.106264	
2	0.292419	0.003101	0.102386	0.001062	0.063179	
3	0.011533	0.002475	0.001515	0.000366	0.000608	



4	0.048458	0.012955	0.006402	0.000080	0.012449
...	...	...	...	...	...
38292	0.002313	0.000599	0.006705	0.000555	0.003945
38293	0.000486	0.012229	0.003478	0.000057	0.003488
38294	0.006975	0.000517	0.008577	0.000178	0.022837
38295	0.999999	0.578736	0.999903	0.008992	0.977605
38296	0.009983	0.002286	0.001205	0.000146	0.002386

	targeted_hate
0	0.001964
1	0.003095
2	0.007714
3	0.002441
4	0.000395
...	...
38292	0.001124
38293	0.001334
38294	0.001180
38295	0.025361
38296	0.000847

[38297 rows x 6 columns]

```
submission_df = pd.concat([tdf.id,y_pred_df],axis=1)
submission_df.to_csv('submission.csv',index = False)
```

## Decision Tree

```
from sklearn.tree import DecisionTreeClassifier
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
X_train = main_df['text']
```

```
y_train = main_df.drop(labels = ['id','text'], axis=1)
```

```
X_test = tdf['text']
```

```
tfidf = TfidfVectorizer()
```

```
x_train_tfidf = tfidf.fit_transform(X_train)
```

```
x_test_tfidf = tfidf.transform(X_test)
```

```
#Fitting Decision Tree classifier to the training set
```

```
classifier= DecisionTreeClassifier()
```

```
classifier.fit(x_train_tfidf, y_train)
```

```
y_pred_dt = classifier.predict_proba(x_test_tfidf)
```

```
#y_pred_lg = classifier.predict_log_proba(x_test_tfidf)
```

```
categories = main_df.iloc[:,2:8].columns
```

```
classifier.get_depth()
```

```

classifier.get_n_leaves()

y_pred_dt
# returning 3D array

y_pred_df_1 = pd.DataFrame(y_pred_dt, columns = categories)
y_pred_df_1

```

### MultiOutput Classifier

```

from sklearn.multioutput import MultiOutputClassifier
from sklearn.svm import LinearSVC
from sklearn.metrics import multilabel_confusion_matrix,
ConfusionMatrixDisplay
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import train_test_split

X_train = main_df['text']

y_train = main_df.drop(labels = ['id', 'text'], axis=1)

X_test = tdf['text']

tfidf = TfidfVectorizer()

x_train_tfidf = tfidf.fit_transform(X_train)
x_test_tfidf = tfidf.transform(X_test)

# Split into training and testing data
X_train, X_test, y_train, y_test = train_test_split(x_train_tfidf,
y_train, test_size=0.33, random_state=42)

# Create the SVM
svm = LinearSVC()

# Make it an Multilabel classifier
multilabel_classifier = MultiOutputClassifier(svm, n_jobs=-1)

# Fit the data to the Multilabel classifier
multilabel_classifier = multilabel_classifier.fit(X_train, y_train)

# Get predictions for test data
y_test_pred = multilabel_classifier.predict(X_test)

# Generate multiclass confusion matrices
matrices = multilabel_confusion_matrix(y_test, y_test_pred)

matrices
array([[26469,   248],
       [   957,  1815]],

      [[29129,    69],

```

```

        [ 225,    66]],

        [[27835,   135],
         [  484,  1035]],

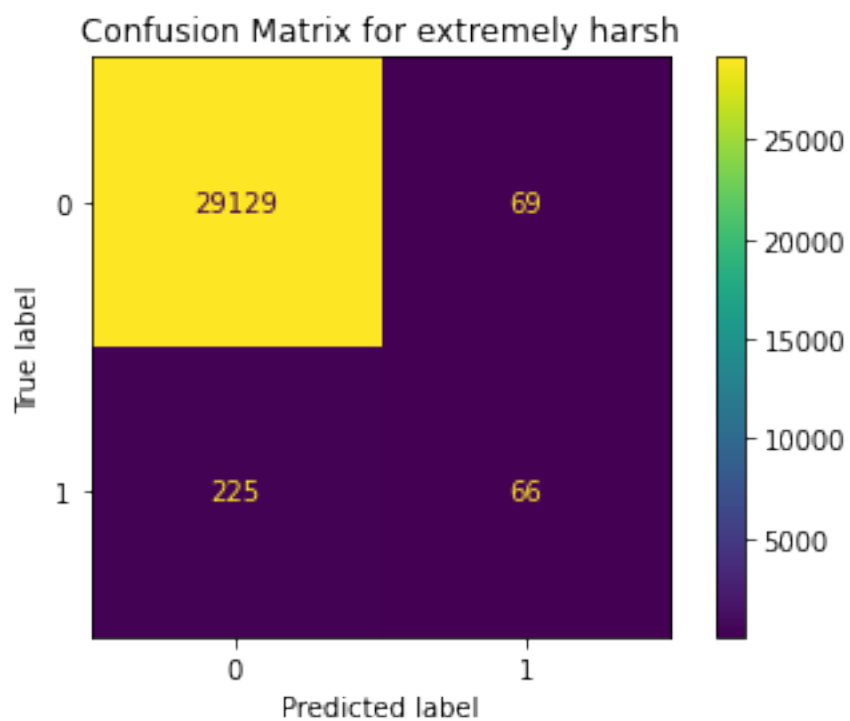

        [[29390,    3],
         [   81,   15]],

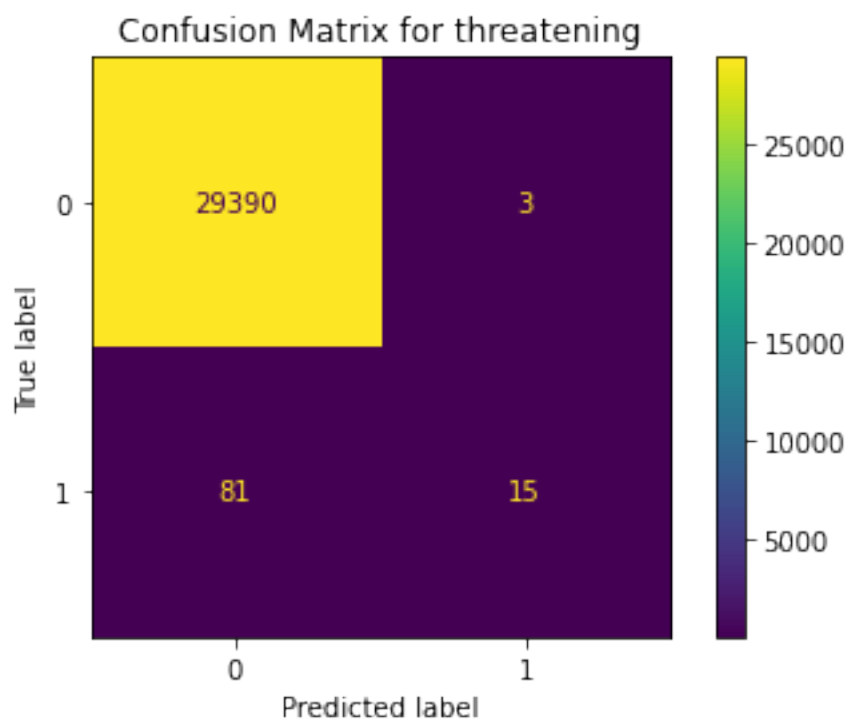

        [[27883,   195],
         [  632,   779]],

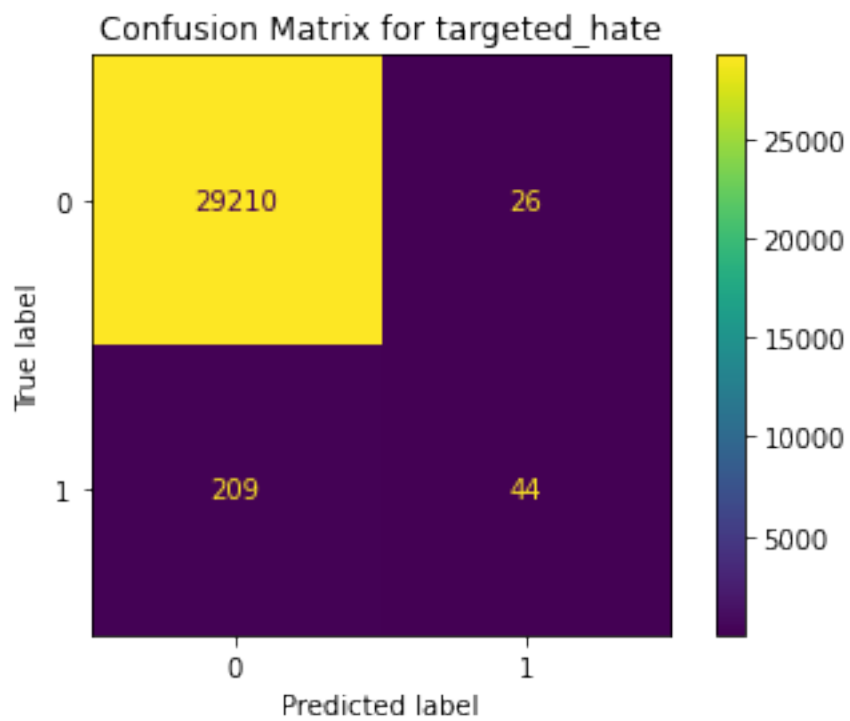

        [[29210,    26],
         [   209,   44]])

cmd = ConfusionMatrixDisplay(matrices[0],
display_labels=np.unique(y_test)).plot()
plt.title('Confusion Matrix for harsh')
plt.show()
cmd = ConfusionMatrixDisplay(matrices[1],
display_labels=np.unique(y_test)).plot()
plt.title('Confusion Matrix for extremely harsh')
plt.show()
cmd = ConfusionMatrixDisplay(matrices[2],
display_labels=np.unique(y_test)).plot()
plt.title('Confusion Matrix for vulgar')
plt.show()
cmd = ConfusionMatrixDisplay(matrices[3],
display_labels=np.unique(y_test)).plot()
plt.title('Confusion Matrix for threatening')
plt.show()
cmd = ConfusionMatrixDisplay(matrices[4],
display_labels=np.unique(y_test)).plot()
plt.title('Confusion Matrix for disrespect')
plt.show()
cmd = ConfusionMatrixDisplay(matrices[5],
display_labels=np.unique(y_test)).plot()
plt.title('Confusion Matrix for targeted_hate')
plt.show()

```







```
y_train.shape
```

```
(59870, 6)
```

```
yy = y_train.head(29489)
```

```
y_test_pred.shape
```

```
(29489, 6)
```

```
# Accuracy
```

```
from sklearn.metrics import accuracy_score
```

```
accuracy_score(yy, y_test_pred)*100
```

```
83.44806538031129
```

## XGBoost

```
from xgboost import XGBClassifier
```

```
from sklearn.feature_extraction.text import TfidfVectorizer
```

```
X_train = main_df['text']
```

```
y_train = main_df.drop(labels = ['id', 'text'], axis=1)
```

```
X_test = tdf['text']
```

```
tfidf = TfidfVectorizer()
```

```
x_train_tfidf = tfidf.fit_transform(X_train)
```

```
x_test_tfidf = tfidf.transform(X_test)
```

```
xgb_clf = XGBClassifier()
```

```
xgb_clf.fit(x_train_tfidf, y_train)
```

```
XGBClassifier(base_score=0.5, booster='gbtree', callbacks=None,  
               colsample_bylevel=1, colsample_bynode=1,  
               colsample_bytree=1,  
               early_stopping_rounds=None, enable_categorical=False,  
               eval_metric=None, gamma=0, gpu_id=-1,  
               grow_policy='depthwise',  
               importance_type=None, interaction_constraints='',  
               learning_rate=0.300000012, max_bin=256,  
               max_cat_to_onehot=4,  
               max_delta_step=0, max_depth=6, max_leaves=0,  
               min_child_weight=1,  
               missing=nan, monotone_constraints='()',  
               n_estimators=100,  
               n_jobs=0, num_parallel_tree=1, predictor='auto',  
               random_state=0,  
               reg_alpha=0, reg_lambda=1, ...)
```

```
y_pred = xgb_clf.predict_proba(x_test_tfidf)
```

```
categories = main_df.iloc[:,2:8].columns
```

```
y_pred_df = pd.DataFrame(y_pred, columns = categories)
```

```
y_pred_df
```

	harsh	extremely_harsh	vulgar	threatening	disrespect	\
0	0.059576	0.003772	0.020237	0.000704	0.024908	
1	0.672233	0.000125	0.006000	0.000057	0.052414	
2	0.153774	0.009824	0.020903	0.003273	0.047614	
3	0.028054	0.000649	0.006546	0.000230	0.004386	
4	0.023089	0.001458	0.007860	0.000088	0.006149	
...	...	...	...	...	...	
38292	0.001813	0.000036	0.000861	0.000020	0.001565	
38293	0.024186	0.003772	0.012379	0.000704	0.014865	
38294	0.018855	0.000219	0.004046	0.000182	0.005519	
38295	0.999288	0.641275	0.978244	0.000736	0.967045	
38296	0.012892	0.000694	0.001294	0.000107	0.003027	

	targeted_hate
0	0.004147
1	0.002133
2	0.007809
3	0.001016
4	0.000375
...	...
38292	0.000081
38293	0.004147
38294	0.000328
38295	0.129180
38296	0.000429

[38297 rows x 6 columns]

```
submission_df = pd.concat([tdf.id,y_pred_df],axis=1)
submission_df.to_csv('submission.csv',index = False)
```

## Ridge Classifier

```
from sklearn.linear_model import Ridge
from sklearn import linear_model
from sklearn.linear_model import RidgeClassifier
from scipy.sparse import hstack

X_train = main_df['text']

y_train = main_df.drop(labels = ['id','text'], axis=1)

X_test = tdf['text']

tfidf = TfidfVectorizer()

x_train_tfidf = tfidf.fit_transform(X_train)
x_test_tfidf = tfidf.transform(X_test)

x_train_tfidf
```



```

<15000x46612 sparse matrix of type '<class 'numpy.float64'>'
  with 395855 stored elements in Compressed Sparse Row format>

x_test_tfidf

<38297x46612 sparse matrix of type '<class 'numpy.float64'>'
  with 901920 stored elements in Compressed Sparse Row format>

predict_dic = pd.DataFrame(columns=['harsh', 'extremely_harsh',
'vulgar', 'threatening', 'disrespect', 'targeted_hate'])

categories = main_df.iloc[:,2:8].columns

char_vectorizer = TfidfVectorizer (
    sublinear_tf = True,
    strip_accents = 'unicode', # Remove accents and perform other
character normalization during the preprocessing step.
                                # 'ascii' is a fast method that only
works on characters that have an direct ASCII mapping.
                                # 'unicode' is a slightly slower method
that works on any characters.
    analyzer = 'char',
    ngram_range = (2, 4), # ngram is the set of n words together.
    min_df = 2,
    max_df = 0.5,
    max_features = 20000
)
char_vectorizer.fit(main_df) # We fit on complete training + test data
so as to achieve a better fit.
train_char_features = char_vectorizer.transform(X_train)
test_char_features = char_vectorizer.transform(X_test)

word_vectorizer = TfidfVectorizer(
    sublinear_tf = True, # It seems unlikely that twenty occurrences
of a term in a document truly carry twenty times the significance of a
single occurrence.
    # Accordingly, there has been considerable research into variants
of term frequency that go beyond counting the number of occurrences of
a term.
    # A common modification is to use instead the logarithm of the
term frequency, which assigns a weight. 1 + log(tf)
    strip_accents = 'unicode', # Remove accents and perform other
character normalization during the preprocessing step.
                                # 'ascii' is a fast method that only
works on characters that have an direct ASCII mapping.
                                # 'unicode' is a slightly slower method
that works on any characters.
    analyzer = 'word', # Whether the feature should be made of word
or character.
    token_pattern = '(?u)\\b\\w\\w+\\b\\w{1}',
    lowercase = False, # Do not convert the uppercase letters into

```

```

lowercase because they carry significance.
stop_words = 'english', # Remove all the stop words of english
# ngram is the set of n words together.
ngram_range = (1, 1), # We consider set of 1 or 2 words together
for tokenization.
min_df = 1,
max_df = 0.25,
norm = 'l2',
max_features = 30000
) #lowercase = true : convert all characters into lower case before
tokenizing
word_vectorizer.fit(main_df) # Apply tfidf fitting on the whole
preprocessed text data so that we achieve a better fitted model.
train_word_features = word_vectorizer.transform(X_train)
test_word_features = word_vectorizer.transform(X_test)

train_features = hstack([train_char_features, train_word_features])
test_features = hstack([test_char_features, test_word_features])

predict_dic['id'] = tdf['id']
overall_score = []
for x in categories:
    print('... Processing {}'.format(x))
    class_column = labels[x].values
    score = 0
    if x in ['targeted_hate', 'threatening']:
        model = LogisticRegression(n_jobs=-1, random_state=0, C=3)
        model.fit(train_features, labels[x])
        predict_dic[x] = model.predict_proba(test_features)
[:,1].tolist()
    else:
        rid = RidgeClassifier(alpha=17, fit_intercept=True,
solver='sag', max_iter=250, random_state=0, tol = 0.0005)
        rid.fit(train_features, labels[x])
        d = rid.decision_function(test_features)
        probs = np.exp(d) / np.sum(np.exp(d))
        predict_dic[x] = probs.tolist()

... Processing harsh
... Processing extremely_harsh
... Processing vulgar
... Processing threatening
... Processing disrespect
... Processing targeted_hate

predict_dic

```

	harsh	extremely_harsh	vulgar	threatening	disrespect	\
0	0.000027	0.000026	0.000026	0.002141	0.000026	
1	0.000022	0.000025	0.000023	0.001209	0.000024	
2	0.000026	0.000026	0.000026	0.002231	0.000026	

3	0.000026	0.000026	0.000026	0.002484	0.000026
4	0.000024	0.000026	0.000025	0.001204	0.000025
...	...	...	...	...	...
38292	0.000024	0.000026	0.000024	0.001290	0.000025
38293	0.000033	0.000027	0.000031	0.006388	0.000030
38294	0.000025	0.000026	0.000026	0.001457	0.000025
38295	0.000029	0.000027	0.000028	0.002641	0.000028
38296	0.000024	0.000026	0.000024	0.001098	0.000025

	targeted_hate	id
0	0.021840	e0ae9d9474a5689a5791
1	0.003734	b64a191301cad4f11287
2	0.007583	5e1953d9ae04bdc66408
3	0.005820	23128f98196c8e8f7b90
4	0.005248	2d3f1254f71472bf2b78
...	...	...
38292	0.007616	64ebe2494b078bc1ec18
38293	0.022400	16259bc32bd803e6acf5
38294	0.004577	1fe631c9625d88a4d492
38295	0.013245	085ab9387dce9d4e0b68
38296	0.008266	4fb0f98b22a4f4469fcf

[38297 rows x 7 columns]

```
submission_df = pd.concat([tdf.id,predict_dic],axis=1)
submission_df.to_csv('submission.csv',index = False)
```